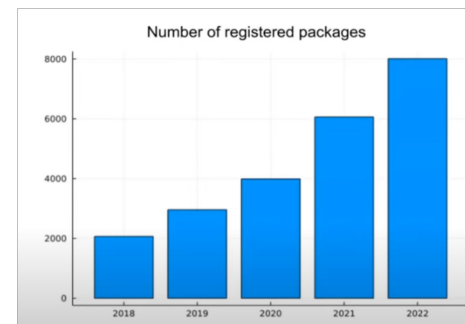
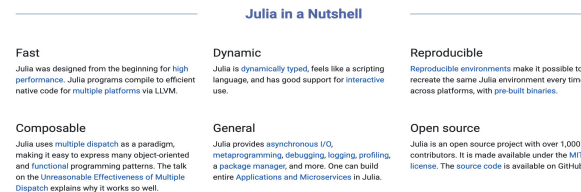



# Introduction to julia

Johannes Blaschke, LBNL, USA

# Julia Brief Walkthrough

- History: started at MIT in the early 2010s (predates Python Numba) <https://julialang.org/blog/2022/02/10years/>
- JuliaHub (formerly Julia Computing) and MIT are major contributors: <https://juliacomputing.com/case-studies>
- First stable release v1.0 in 2018, **v1.11 as of 2024** <https://julialang.org/>
- Open-source GitHub-hosted packages and ecosystem with MIT permissive license: <https://github.com/JuliaLang/julia>
- Community: annual JuliaCon summer conference: <https://juliacon.org/2024/>



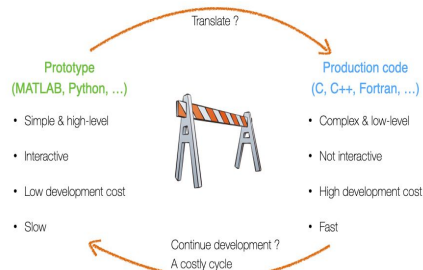
95% of Julia packages in the registry had some form of CI ([youtube.com/watch?v=9YWwiFbaRx8](https://youtube.com/watch?v=9YWwiFbaRx8))

# Julia's value proposition for science


- Designed for “scientific computing” (Fortran-like) and “data science” (Python-like) with **performant kernel code via LLVM compilation**
- Lightweight **interoperability with existing Fortran and C libraries**
- Julia is a **unifying workflow language** with a **coordinated ecosystem**

“Julia **does not** replace Python, but the costly workflow process around **Fortran+Python+X, C+X, Python+X** or **Fortran+X** (e.g. GPUs, simulation + data analysis)”


**where X = { conda, pip, pybind11, cython, Python, C, Fortran, C++, OpenMP, OpenACC, CUDA, HIP, CMake, numpy, scipy, matplotlib, Jupyter, ...}**



<https://pde-on-gpu.vaw.ethz.ch/lecture7>




<https://developer.nvidia.com/blog/gpu-computing-julia-programming-language/>



<https://juliadatascience.io/>

Rich data science ecosystem

<https://quantumzeitgeist.com/learning-the-julia-programming-language-for-free/>



# Julia and HPC



# Why Julia for HPC?

~~Walks like Python~~, talks like Lisp, runs like Fortran

HPC suffers from the many language problem;  
Domain experts and performance engineers use  
different programming languages:

Communication and Collaboration bottleneck

<https://www.nature.com/articles/d41586-019-02310-3>

**Julia: come for the syntax,  
stay for the speed**



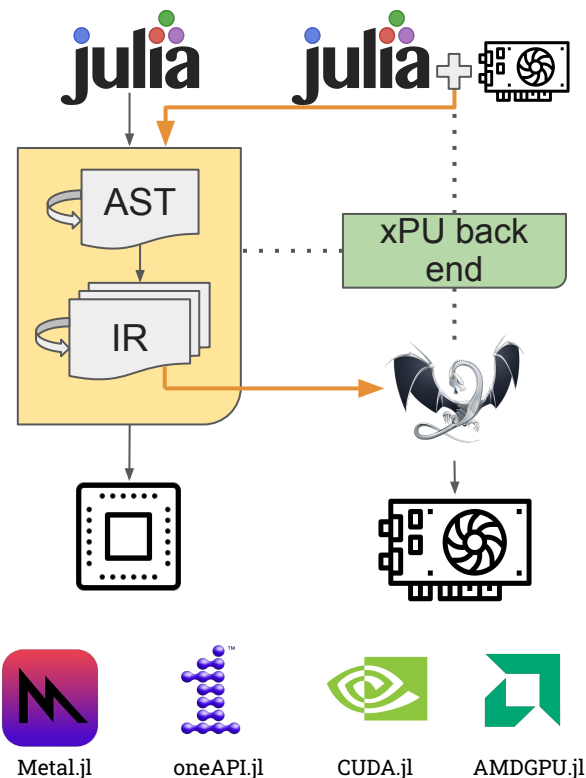
Katie Hyatt  
@kslimes



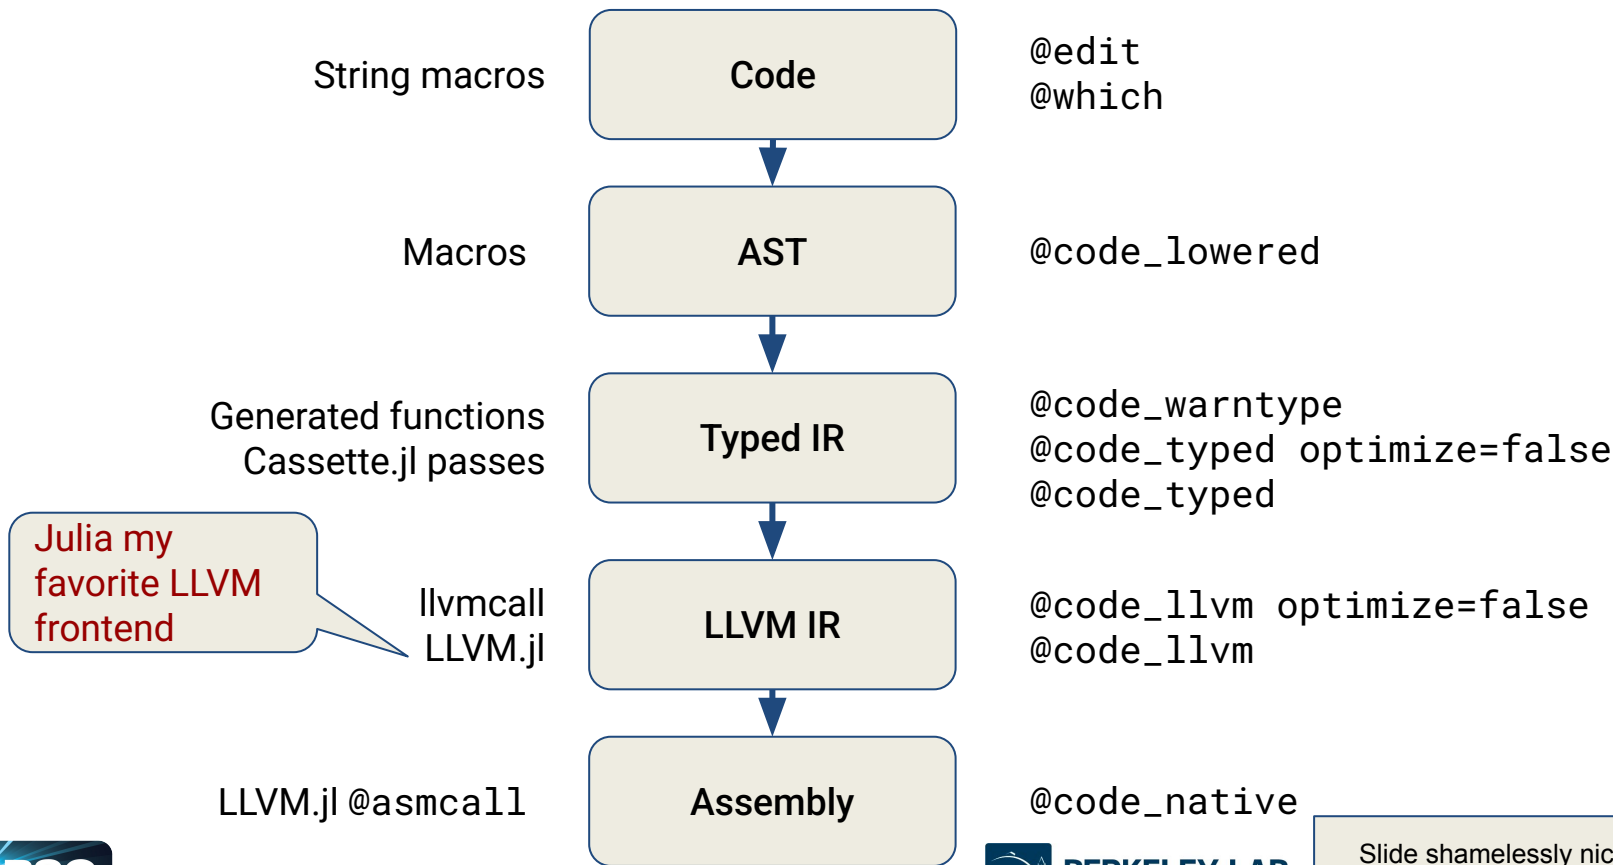
My port of our research code from CPU-based C++ to GPU-accelerated [#julia](#) is so much faster (1 week -> 1 hour walltimes) and so much easier to add stuff to... it's a nice holiday gift to myself ❄️☀️.



## Rich GPU Ecosystem



# Introspection and staged metaprogramming



# Multiple Dispatch => Situation-Dependent Code Generation

- Let's take a closer look at Multiple Dispatch:
  - A function's implementation (method) is selected depending on the input types of the

```
In [40]: function double_int(x::Int)
          return 2*x
        end

function double_int(x::AbstractFloat)
    y = floor{Int}(x)
    r = x - y
    return 2*y + r
end

Out[40]: double_int (generic function with 2 methods)
```

```
In [30]: double_int(10)
```

```
Out[30]: 20
```

```
In [31]: double_int(10.1)
```

```
Out[31]: 20.1
```

We can list the methods for a function using the `methods` function:

```
In [32]: methods(double_int)
```

```
Out[32]: # 2 methods for generic function double_int:
          • double_int(x::AbstractFloat) in Main at In[29]:5
          • double_int(x::Int64) in Main at In[29]:1
```

# Multiple Dispatch => Situation-Dependent Code Generation

- Let's take a closer look at Multiple Dispatch:
  - A function's implementation (method) is selected depending on the input types of the

```
In [40]: function double_int(x::Int)
          return 2*x
        end

function double_int(x::AbstractFloat)
    y = floor{Int}(x)
    r = x - y
    return 2*y + r
end

Out[40]: double_int (generic function with 2 methods)
```

The `@code_lowered` macro gives a (still somewhat abstract) idea what Julia actually does.

```
In [33]: @code_lowered double_int(2)
```

```
Out[33]: CodeInfo(
  1 - %1 = 2 * x
      └─ return %1
)
```

This picks up the method for `x` as an integer, and similarly we can see what Julia does when `x` is a float:

```
In [34]: @code_lowered double_int(2.1)
```

```
Out[34]: CodeInfo(
  1 -      y = Main.floor(Main.Int, x)
      r = x - y
      %3 = 2 * y
      %4 = %3 + r
      └─ return %4
)
```



# Multiple Dispatch => Situation-Dependent Code Generation

- Let's take a closer look at Multiple Dispatch:
  - Following dispatch, Typed IR is translated to LLVM IR

```
In [40]: function double_int(x::Int)
          return 2*x
        end

function double_int(x::AbstractFloat)
    y = floor{Int}(x)
    r = x - y
    return 2*y + r
end
```

Out[40]: double\_int (generic function with 2 methods)

## Julia is a REPL for LLVM:

```
In [41]: @code_llvm double_int(2)

; @ In[40]:1 within `double_int`
define i64 @julia_double_int_2028(i64 signext %0) #0 {
top:
; @ In[40]:2 within `double_int`
; @ int.jl:88 within `*`
    %1 = shl i64 %0, 1
; L
    ret i64 %1
}
```

# Magic of Julia

Abstraction, Specialization, and Multiple Dispatch

## 1. **Abstraction** to obtain generic behavior:

Encode behavior in the type domain:

```
transpose(A::Matrix{Float64})::Transpose{Float64,Matrix{Float64}}
```

Did I really need to move memory for that transpose?

## 2. **Specialization** of functions to produce optimal code

## 3. **Multiple-dispatch** to select optimized behavior

```
rand(N, M) * rand(K, M)'  
Matrix * Transpose{Matrix}
```

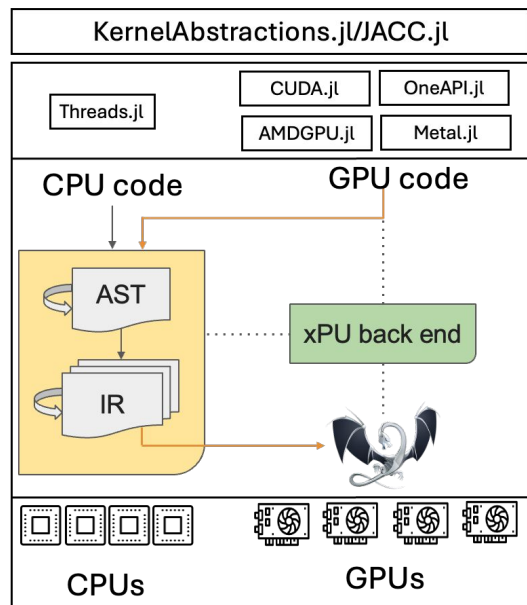
compiles to

```
function mul!(C::Matrix{T}, A::Matrix{T}, tB::Transpose{<:Matrix{T}}, a, b) where {T<:BlasFloat}  
    gemm_wrapper!(C, 'N', 'T', A, B, MulAddMul(a, b))  
end
```

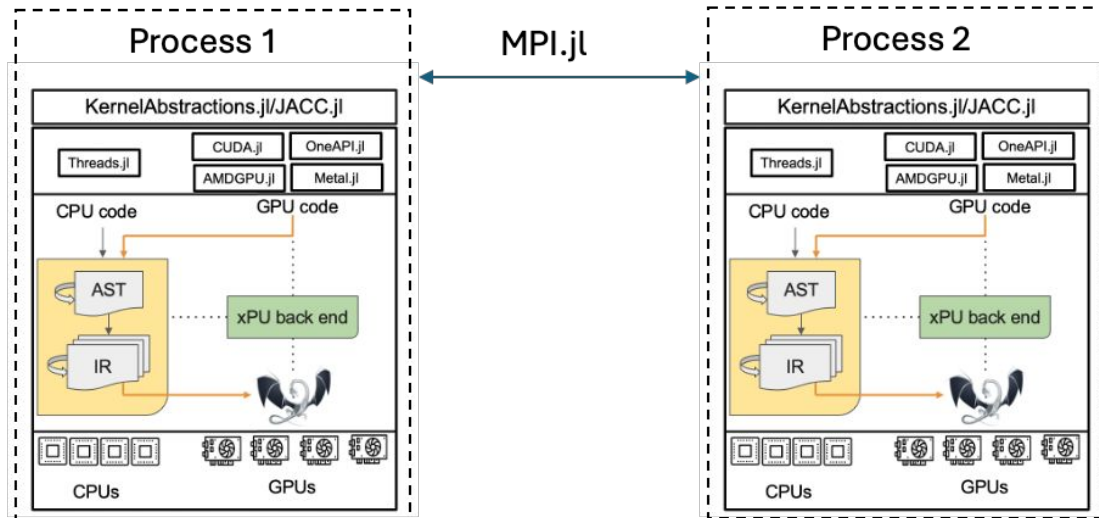
No I did not! I know  $AB^T$  is the dot product of every row of A with every row of B.

# Julia Parallel Paradigm

## Shared-memory



## Distributed-memory



# Array programming

**using** LinearAlgebra

```
loss(w,b,x,y) = sum(abs2, y - (w*x .+ b)) / size(y,2)
loss∇w(w, b, x, y) = ...
lossdb(w, b, x, y) = ...
```

```
function train(w, b, x, y ; lr=.1)
    w -= lmul!(lr, loss∇w(w, b, x, y))
    b -= lr * lossdb(w, b, x, y)
    return w, b
end
```

```
n = 100; p = 10
x = randn(n,p)'
y = sum(x[1:5,:]; dims=1) .+ randn(n)'*0.1
w = 0.0001*randn(1,p)
b = 0.0
```

```
for i=1:50
    w, b = train(w, b, x, y)
end
```


```
x = CuArray(x)
y = CuArray(y)
w = CuArray(w)
```



Rapid software prototyping for heterogeneous and distributed platforms

Besard T., Churavy V., Edelman A., De Sutter B.

([doi:10.1016/j.advenzsoft.2019.02.002](https://doi.org/10.1016/j.advenzsoft.2019.02.002))



# KernelAbstractions.jl

- **GPU-centric** programming model
- SIMT execution model
- Multiple GPU backends
- Implementation using **macros** and **method overlay tables**
- GPU backends:
  - CUDA, AMDGPU, oneAPI, Metal
  - 172 - 188 LOC
- CPU execution as a fallback and for debugging
- Widely adopted in the JuliaHPC and JuliaGPU ecosystem

```
@kernel function lmem_copy_kernel!(output, @Const(input))  
    I, J = @index(Global, NTuple)  
    i, j = @index(Local, NTuple)  
    N = @uniform @groupsize()[1]  
    M = @uniform @groupsize()[2]  
    tile = @localmem eltype(output) (N, M)  
  
    @inbounds tile[i, j] = input[I, J]  
    @synchronize  
    @inbounds output[I, J] = tile[i, j]  
end
```

# (advanced) LLVM + Julia

Julia provides  
interfaces to the  
LLVM backend.

Eg.:


- `loopinfo`
- `llvmcall`

```
[16]: macro unroll(expr)
      expr = loopinfo("@unroll", expr, (Symbol("llvm.loop.unroll.full"),))
      return esc(expr)
    end

    for (jlf, f) in zip((:+, :*, :-), (:add, :mul, :sub))
      for (T, llvmT) in ((:Float32, "float"), (:Float64, "double"))
        ir = ""
        %x = f$f contract nsz $llvmT %0, %1
        ret $llvmT %x
        ""
        @eval begin
          # the @pure is necessary so that we can constant propagate.
          @inline Base.@pure function $jlf(a::$T, b::$T)
            Base.llvmcall($ir, $T, Tuple{$T, $T}, a, b)
          end
        end
      end
    end
    @eval function $jlf(args...)
      Base.$jlf(args...)
    end
  end
```

# In a Nutshell...Why Julia??


- *Think in Python/Fortran, but now imagine that it works “well” on HPC*
- A JIT “technical” language on top of LLVM
  - Easy-to-use and agile interface with expected performance
- Julia syntax is for Science and is Open Source
- Native syntax (no manual implemented wrappers of wrapper of ...) for HPC support, e.g. [www.juliagpu.org](http://www.juliagpu.org)
  - Threads, CUDA, AMDGPU, OneAPI, MPI, DAGGER, etc.
- Interoperability: C and Fortran, support for AI: FluxML
- Integrated and efficient support for packaging, reproducibility, CI/CD, ...
- All this makes Julia’s ecosystem and community motivated by performance and productivity -> e.g. **rapid prototyping, “throw-away” code, scientific CPU/GPU access, HPC stack**



High-performance GPU programming in a high-level language.

# Why NOT Julia??

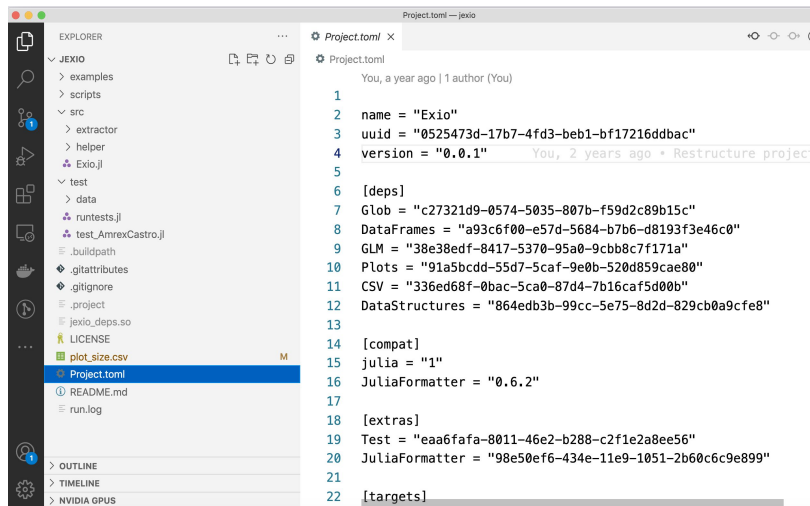
- Existing large investments in other languages
- Long-term ROI: Package support? Stable not standard
- ~~Ecosystem is not mature: Tooling? (HPC)~~
- Out of scope for my application needs
- I'm simply more comfortable in another language





# Julia Brief Walkthrough

- Reproducibility is in the core of the language:
  - Interactive: Jupyter, Pluto.jl
  - Packaging Pkg.jl
  - Environment Project.toml
  - Testing Test.jl
- Just-in-time or Ahead-of-time compilation with PackageCompiler.jl
- Powerful metaprogramming for code instrumentation: @profile, @time, @testset, @test, @code\_lvm, @code\_native, @inbounds,
- Interoperability is key: @ccall, @cxx, PyCall, CxxWrap.jl



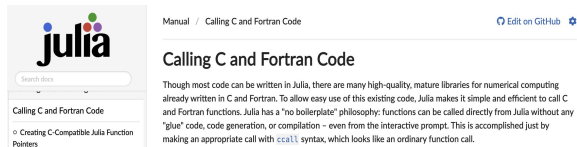
```
Project.toml
You, a year ago | 1 author (You)

1
2 name = "Exio"
3 uuid = "0525473d-17b7-4fd3-beb1-bf17216ddbdc"
4 version = "0.0.1"
5
6 [deps]
7 Glob = "c27321d9-0574-5035-807b-f59d2c89b15c"
8 DataFrames = "a93c6f00-e57d-5684-b7b6-d8193f3e46c0"
9 GLM = "38e3edf-8417-5370-95a0-9cbb8c7f171a"
10 Plots = "91a5bcd-55d7-5caf-9e0b-520d859cae80"
11 CSV = "336ed68f-0bac-5ca0-87d4-7b16caf5d00b"
12 DataStructures = "864edb3b-99cc-5e75-8d2d-829cb0a9cfe8"
13
14 [compat]
15 julia = "1"
16 JuliaFormatter = "0.6.2"
17
18 [extras]
19 Test = "eaa6fafa-8011-46e2-b288-c2f1e2a8ee56"
20 JuliaFormatter = "98e50ef6-434e-11e9-1051-2b60c6c9e899"
21
22 [targets]
```



```
runtests.jl
You, 2 years ago | 1 author (You)

1
2 using Test, Base.Filesystem
3
4 import Exio
5
6 @testset "test_AmrexCastro" begin
7     include("test_AmrexCastro.jl")
8 end;
9
10 @testset "test_Exio.input_parser_docstring" begin
11     @test println(@doc Exio._input_parser) == nothing
12 end;
13
```



<https://github.com/ornl-training/julia-basics>